

Cloud-Based Document Management System with AWS Elasticsearch

*A Course Project Report Submitted in partial fulfillment of the course
requirements for the award of grades in the subject of*

**CLOUD BASED AIML SPECIALITY
(22SDCS07A)**

by

**Deekshitha Bethireddy
(2210030070)**

Under the esteemed guidance of

**Ms. P. Sree Lakshmi
Assistant Professor,
Department of Computer Science and Engineering**



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

K L Deemed to be UNIVERSITY

*Aziznagar, Moinabad, Hyderabad,
Telangana, Pincode: 500075*

April 2025

K L Deemed to be UNIVERSITY

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING



Certificate

This is Certified that the project entitled “**Cloud-Based Document Management System with AWS Elasticsearch**” which is an experimental & Simulation work carried out by Deekshitha Bethireddy (2210030070), in partial fulfillment of the course requirements for the award of grades in the subject of **CLOUD BASED AIML SPECIALITY**, during the year **2024-2025**. The project has been approved as it satisfies the academic requirements.

Ms.P.Sree Lakshmi

Course Coordinator

Dr. Arpita Gupta

Head of the Department

Ms. P. Sree Lakshmi

Course Instructor

CONTENTS

Page No.

1. Introduction	1
2. AWS Services Used as part of the project	2
3. Steps involved in solving project problem statement	4
4. Stepwise Screenshots with brief description	6
5. Learning Outcomes	15
6. Conclusion	16
7. References	17

1. INTRODUCTION

In the digital era, organizations generate vast amounts of documents, necessitating efficient storage, retrieval, and search capabilities. Traditional document management systems often face challenges with scalability, search performance, and infrastructure management. A serverless cloud-based solution addresses these issues by leveraging managed services for scalability, cost-efficiency, and performance. The Cloud-Based Document Management System with AWS Elasticsearch utilizes Amazon OpenSearch Service (formerly AWS Elasticsearch) to enable fast, full-text search and analytics on documents stored in Amazon S3, integrated with serverless AWS services for automation and accessibility [1][2].

This project builds a scalable, serverless document management system that allows users to upload documents to an S3 bucket, automatically index their content in OpenSearch, and perform keyword-based searches via an API Gateway. The system supports various file types, such as text (e.g., “design_principles.txt”), and uses AWS Lambda for automating text extraction and indexing. It is particularly valuable for applications like digital libraries, enterprise content management, and archival systems, where quick document retrieval enhances productivity and decision-making [3][4].

By integrating services like Amazon S3, AWS Lambda, Amazon OpenSearch Service, and API Gateway, this project delivers a fault-tolerant, scalable, and secure document management solution. The use of serverless technologies minimizes operational overhead, while OpenSearch ensures robust search capabilities. This project demonstrates the power of AWS-native services in building modern, cloud-based document management systems that meet the demands of dynamic business environments [5].

This architecture is particularly valuable for industries such as education, healthcare, and corporate environments, where quick access to documents improves decision-making and operational efficiency.

2. AWS Services Used as part of the project

1. Amazon S3 – Scalable Document Storage

Amazon S3 serves as the primary storage layer for documents, offering:

- Highly durable and scalable storage for files like PDFs and text documents.
- Event notifications to trigger AWS Lambda functions upon file uploads.
- Integration with other AWS services for processing and archival [5].
- Supports lifecycle policies for cost-effective archival of older documents.



2. Amazon OpenSearch Service – Full-Text Search and Analytics

Amazon OpenSearch Service (formerly AWS Elasticsearch) powers the search functionality, enabling:

- Full-text indexing and searching of document content using a documents index.
- Real-time search queries for keywords (e.g., “aboutMe” or “design”).
- Scalable and managed search infrastructure for high-performance retrieval [1].
- Scales automatically to handle large document collections.



3. AWS Lambda – Serverless Compute for Automation

AWS Lambda automates document processing and indexing, providing:

- Text extraction from uploaded documents using libraries like pymupdf.
- Indexing extracted content into OpenSearch for searchability.

- Event-driven processing triggered by S3 uploads or API Gateway requests [5].



4. AWS API Gateway – Secure API for Search Queries

AWS API Gateway enables secure and scalable access to the document search functionality:

- Exposes RESTful endpoints for querying the OpenSearch index.
- Integrates with Lambda to process search requests and return results.
- Implements authentication and throttling to ensure secure access.
- Supports testing via tools like Postman for debugging and validation [6].



5. AWS IAM (Identity & Access Management) – Security & Access Control

AWS IAM ensures secure access to AWS resources:

- Role-based access control (RBAC) for users like dms-admin and roles like dms-ExtractText-role-gvkjhfb.
 - Policies granting permissions for S3, OpenSearch, Lambda, and API Gateway.
 - Encryption and authentication to protect document data [5].
 - Tracks access logs via CloudTrail for auditing purposes.
- This enables monitoring and review of all system activities.



6. Amazon CloudWatch – Monitoring & Logging

Amazon CloudWatch monitors system performance and logs:

- CloudWatch tracks AWS Lambda executions, S3 events, and API Gateway performance metrics in real time.
- CloudWatch triggers alerts for operational anomalies such as failed uploads, high error rates, or service downtime.
- Tracks Lambda function executions, API Gateway requests, and OpenSearch performance.
- Logs errors like NoSuchKey or 403 Forbidden for debugging.
- Provides metrics for optimizing resource utilization [5].



3. STEPS INVOLVED IN SOLVING PROJECT PROBLEM STATEMENT

Understanding the Problem Statement

- Identify the need for a scalable document management system to store, index, and search documents like “about_design_v2.txt”.
- Define requirements for full-text search, automation of indexing, and secure access via APIs.
- Ensure scalability, low latency, and cost-efficiency using serverless AWS services.

Designing the AWS Architecture

- Use Amazon S3 as the storage layer for documents, with event triggers for new uploads [5].
- Implement Amazon OpenSearch Service to index document content in a documents index for fast searches [1].
- Deploy AWS Lambda to extract text from documents (e.g., using pymupdf for PDFs) and index it in OpenSearch [7].
- Configure AWS API Gateway to provide a RESTful interface for searching documents by keywords [6].
- Use Amazon CloudWatch to monitor system performance and log errors [5].

Data Ingestion & Processing

- Upload documents (e.g., “AI-ML_internship.pdf”) to an S3 bucket (documentt-storage-bucket).
- Trigger a Lambda function on S3 upload to extract text using pymupdf or similar libraries.
- Index extracted text in the OpenSearch documents index for searchability.
- Handle errors like NoSuchKey by verifying S3 object keys and paths.

Security & Access Control

- Set up IAM roles (e.g., dms-ExtractText-role-gvkjhfb) with policies like AmazonS3FullAccess and AmazonOpenSearchServiceFullAccess [5].
- Resolve 403 Forbidden errors for user dms-admin by updating OpenSearch domain policies and role mappings [8].
- Implement encryption for data at rest and in transit.

Monitoring & Optimization

- Use CloudWatch to monitor Lambda logs, API Gateway metrics, and OpenSearch performance [5].
- Optimize Lambda functions by using Lambda Layers for dependencies like pymupdf to avoid compatibility issues.
- Adjust OpenSearch shard configurations for efficient indexing and search.

Deployment & Testing

- Deploy the system, including S3 bucket, OpenSearch domain, Lambda functions, and API Gateway.
- Test search functionality using Postman with queries like “design” or “internship”.
- Debug issues like DNS errors (getaddrinfo ENOTFOUND) or empty search results by verifying configurations.

Search & Retrieval

- Query the OpenSearch index via API Gateway to retrieve documents matching keywords.
- Validate search results for accuracy and performance using tools like Postman.

4. STEPWISE SCREENSHOTS WITH BRIEF DESCRIPTION

Step 1: Create an IAM User for Secure Access

- Navigate to IAM, create a user (dms-admin), and attach policies:
 - AmazonS3FullAccess
 - AmazonOpenSearchServiceFullAccess
 - AWSLambdaFullAccess
- Generate an access key for programmatic access.

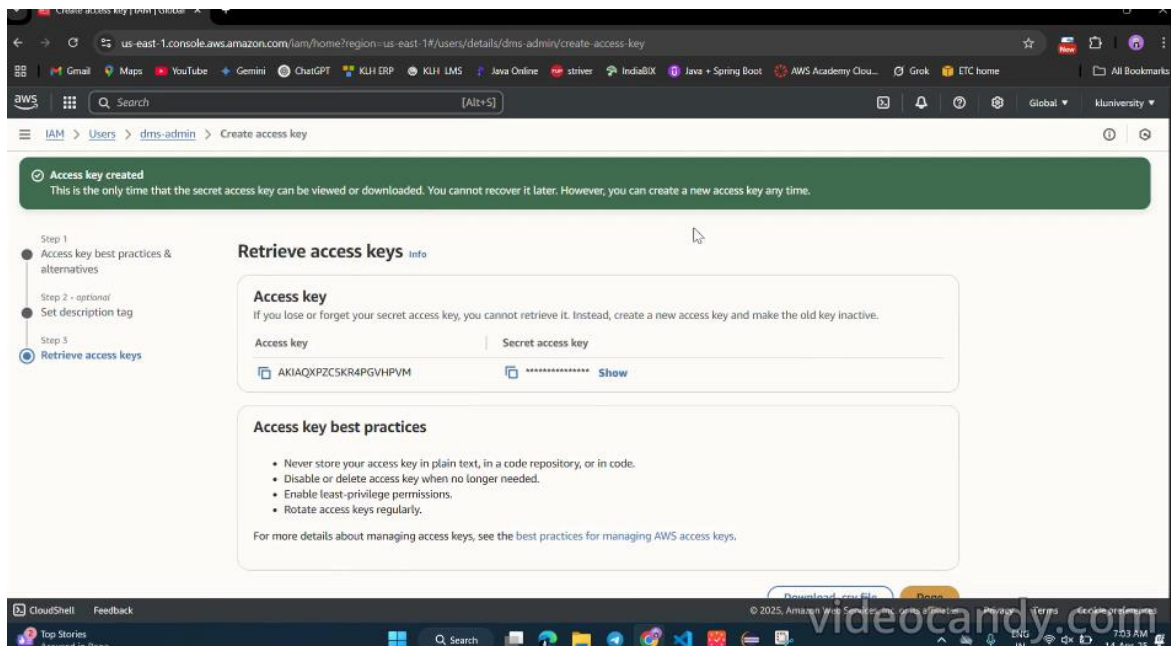


Fig 4.1.1 Creating the IAM user dms-admin and access key.

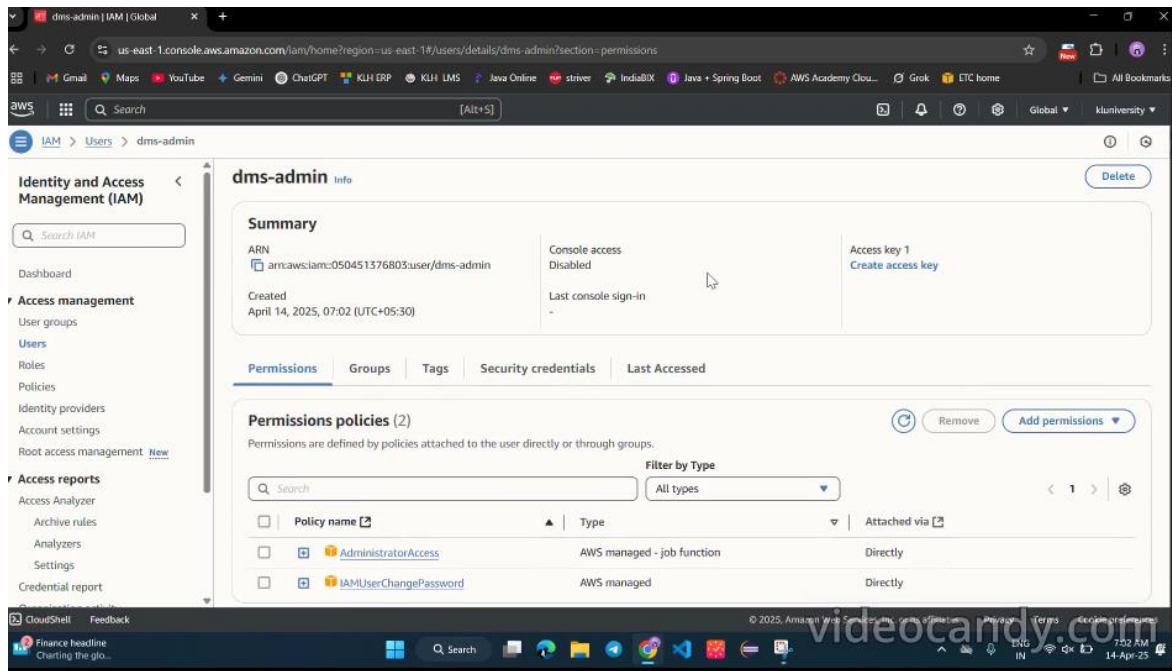


Fig 4.1.2 Attaching required policies for S3, OpenSearch, and Lambda.

Step 2 : Create an S3 Bucket for Document Storage

- Create a general-purpose S3 bucket (documentt-storage-bucket) for storing documents.
- Configure event notifications to trigger Lambda on file uploads.

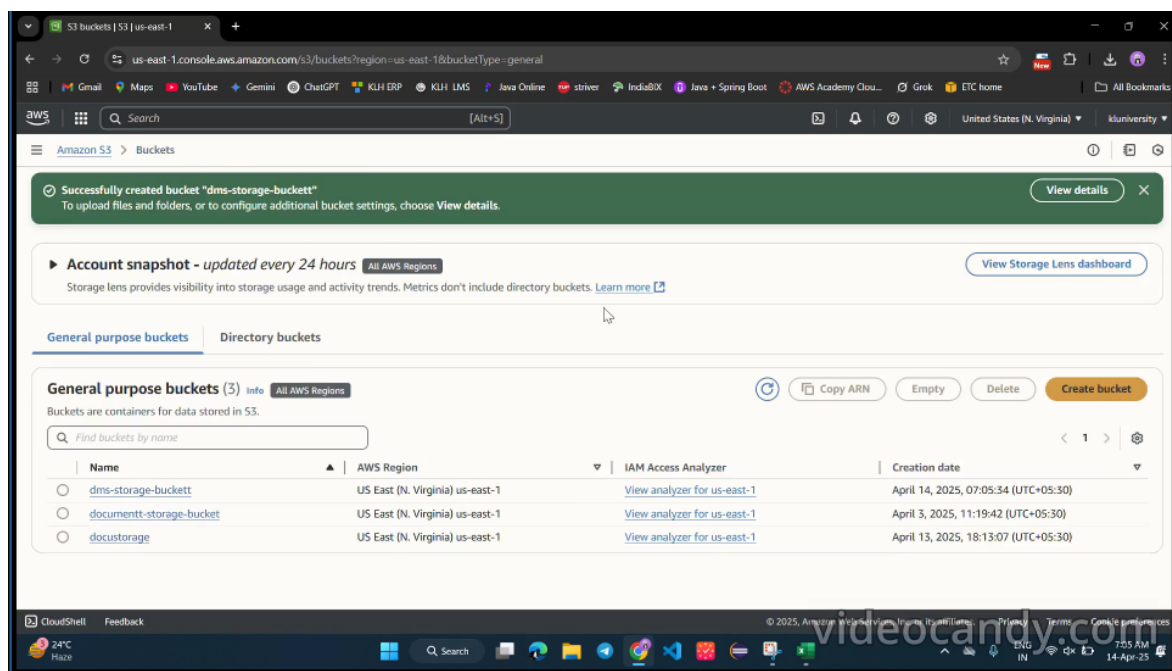


Fig 4.2.1 Successfully created S3 bucket.

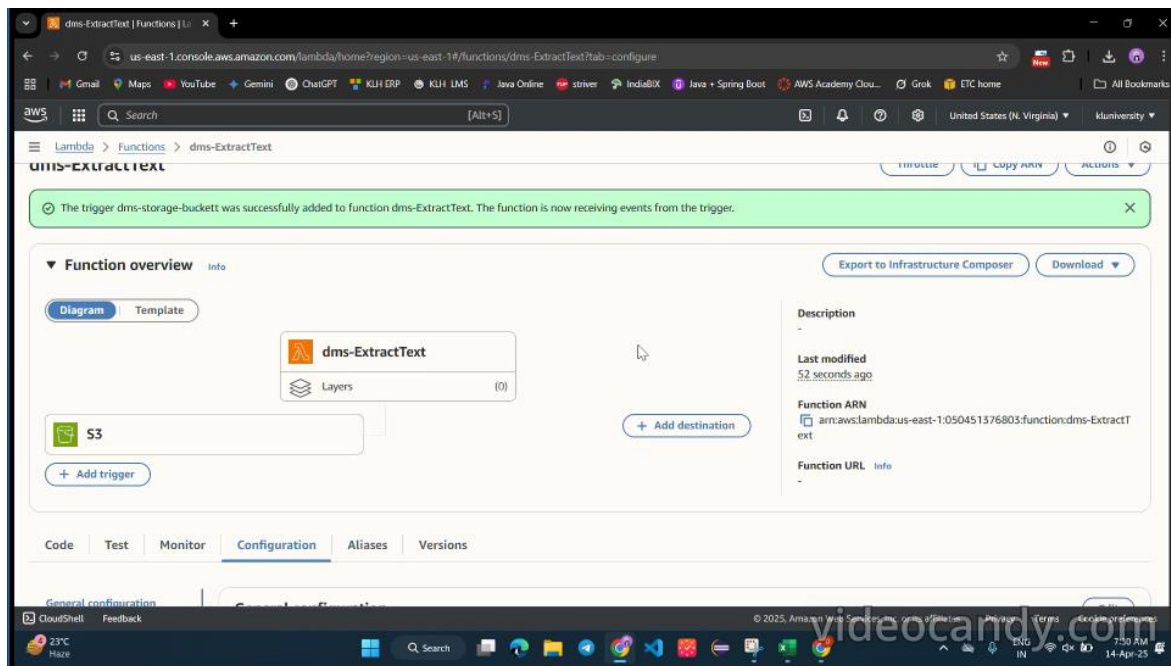


Fig 4.2.2 Configuring S3 event notification for Lambda.

Step 3 : Set Up Amazon OpenSearch Service

- Create an OpenSearch domain and configure the documents index for document content.
- Click “Create domain” and name the domain (e.g., dms-search).
- Select the “Development and testing” deployment type for cost-efficiency during prototyping.
- Choose the latest OpenSearch version available (e.g., OpenSearch 2.11) for optimal features.
- Configure instance types (e.g., t3.small.search) and set 1 data node for simplicity.
- Verify domain status as active.

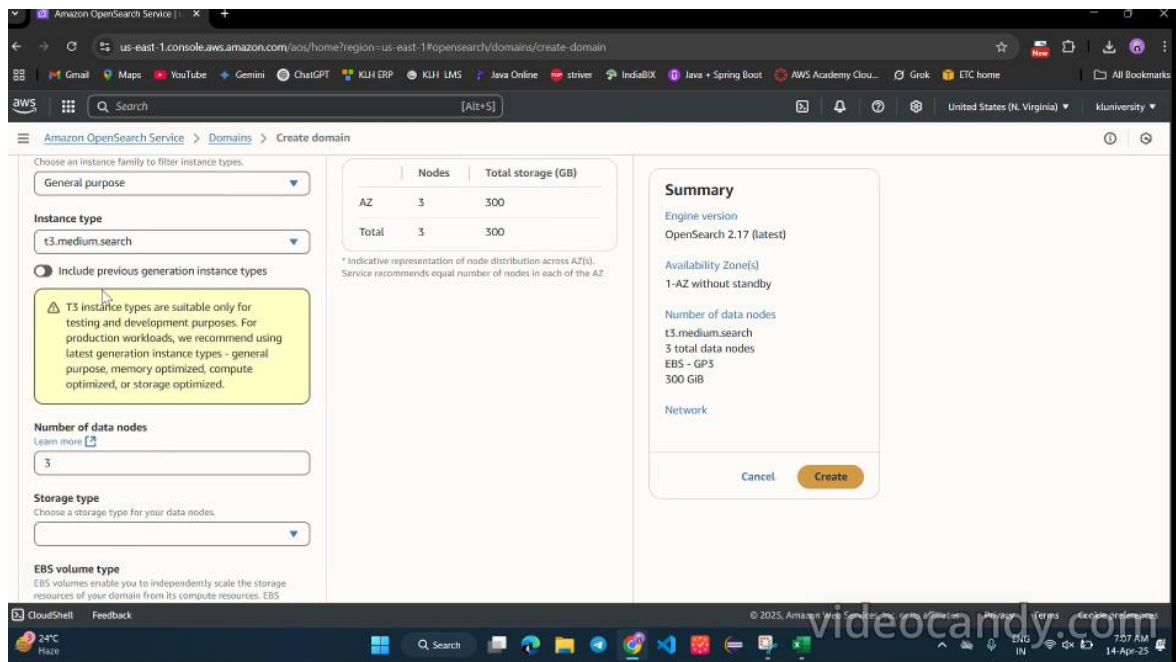


Fig 4.3.1 Accessing the OpenSearch Service console.

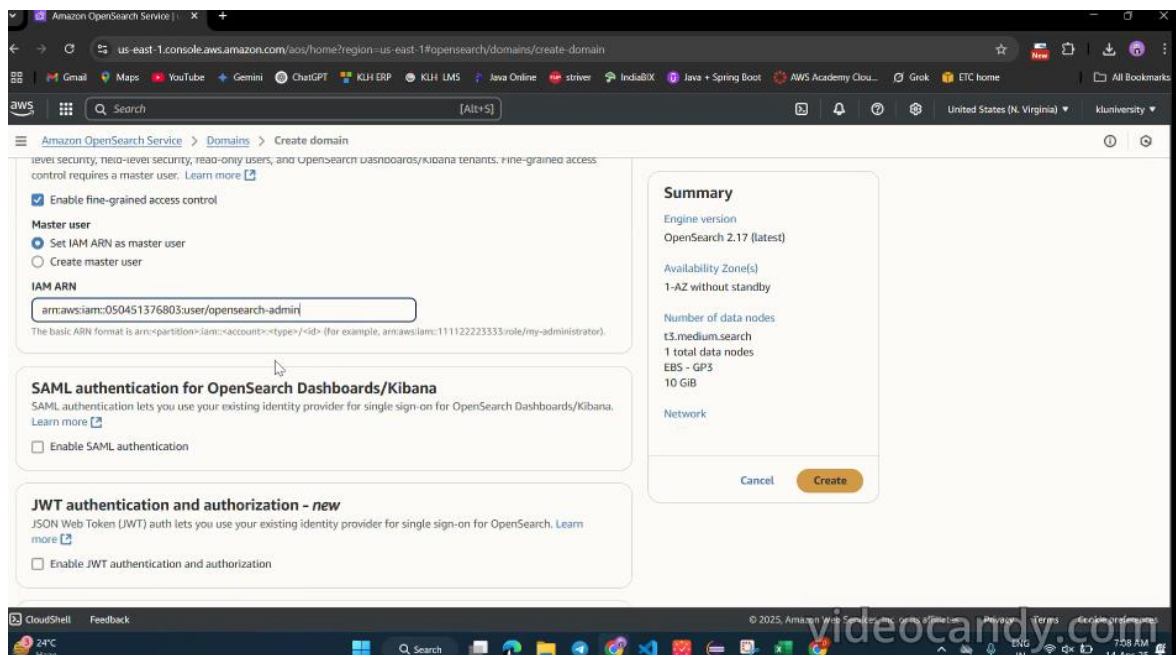


Fig 4.3.2 Configuring the OpenSearch domain settings.

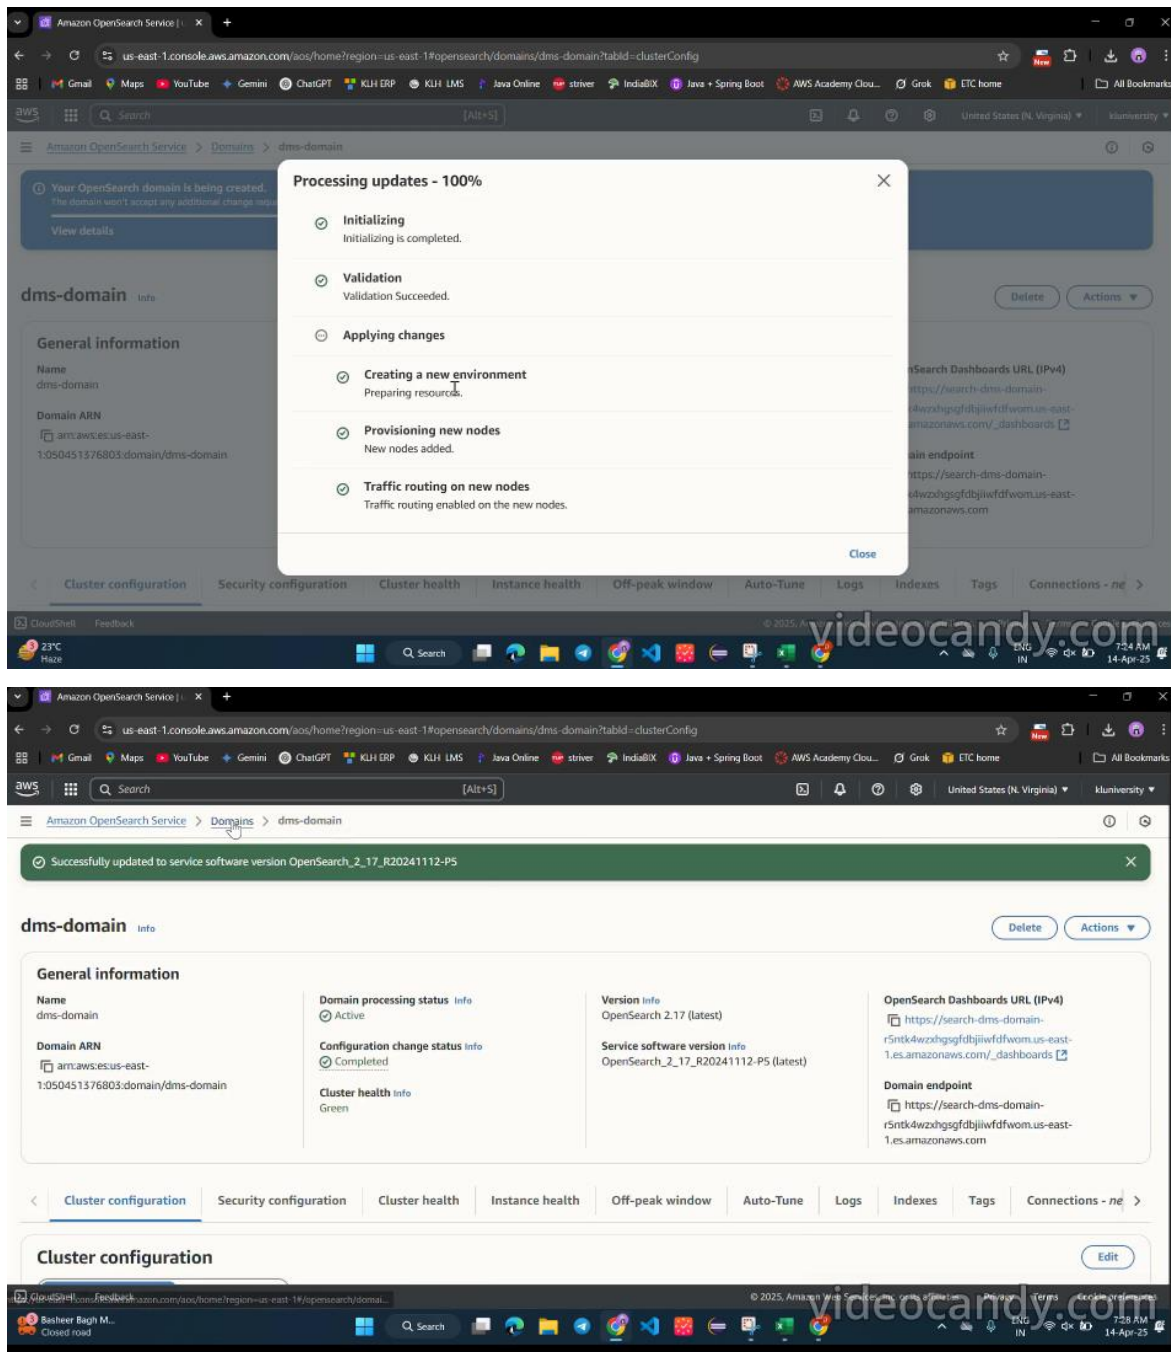


Fig 4.3.3 OpenSearch domain status successful

Step 4 : Create a Lambda Function for Indexing

- Create a Lambda function with an IAM role (dms-ExtractText-role-gvckjfb).
- Add a Lambda Layer for pymupdf to handle PDF text extraction.
- Write code to extract text from S3 uploads and index it in OpenSearch.

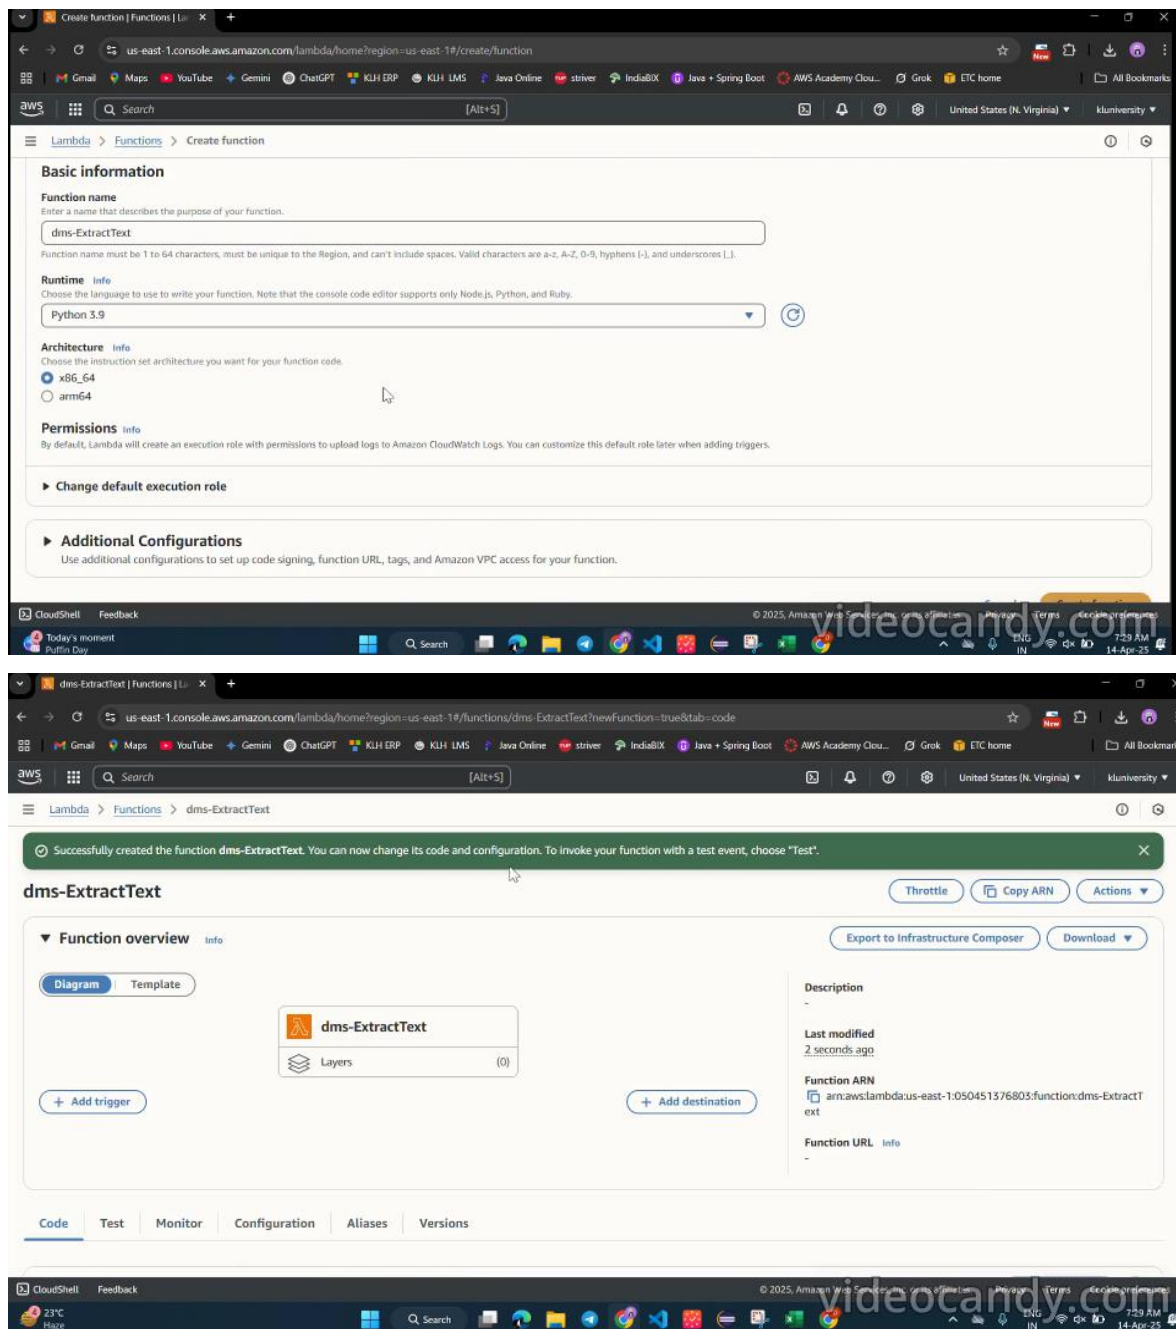


Fig 4.4.1 Creating the Lambda function.

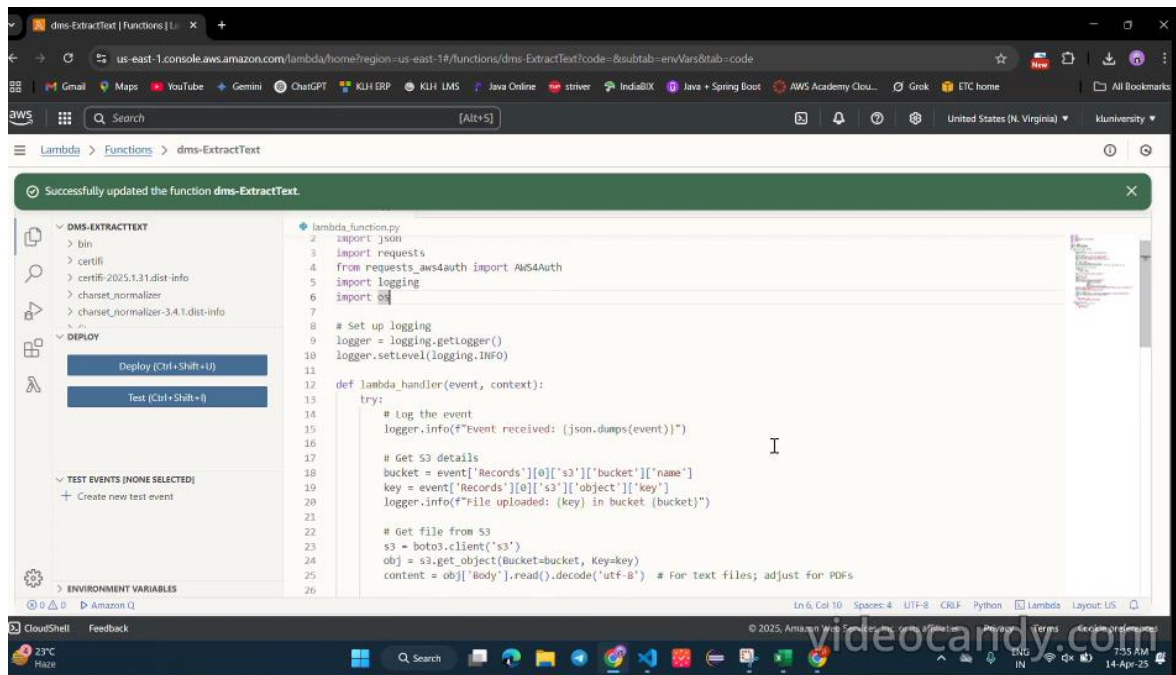


Fig 4.4.2 Lambda function code for indexing.

Step 5 : Upload and Test Documents

- Upload a test file (about_design_v2.txt) to the S3 bucket.
- Verify that the Lambda function indexes the file content in OpenSearch.

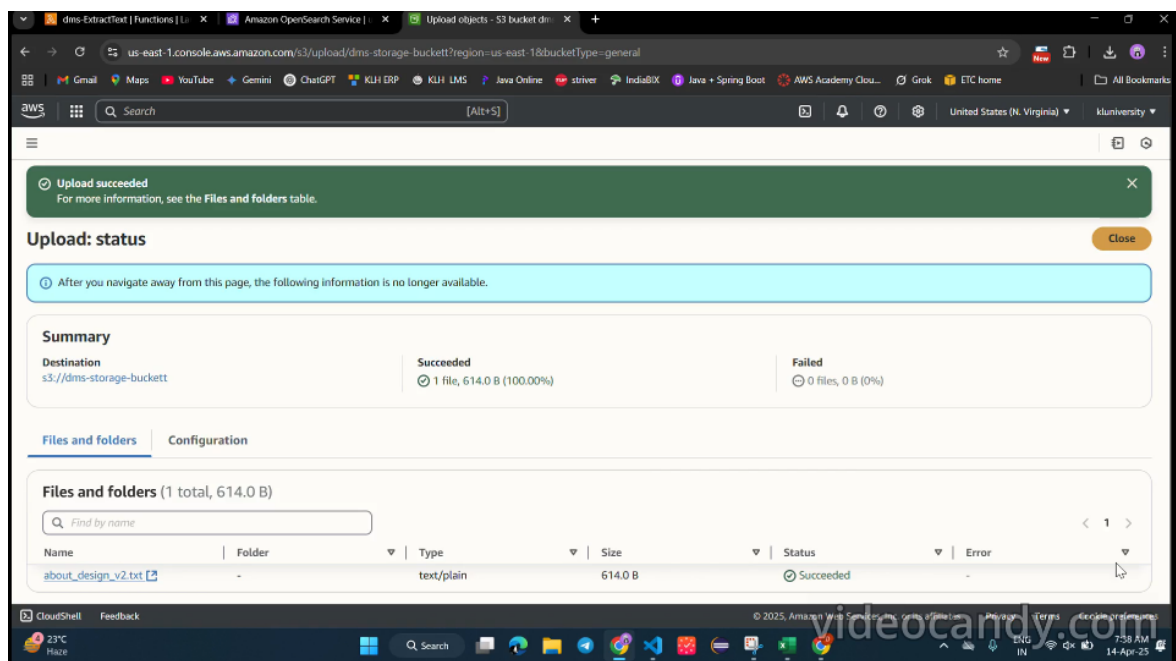


Fig 4.5.1 Uploading about_design_v2.txt to S3.

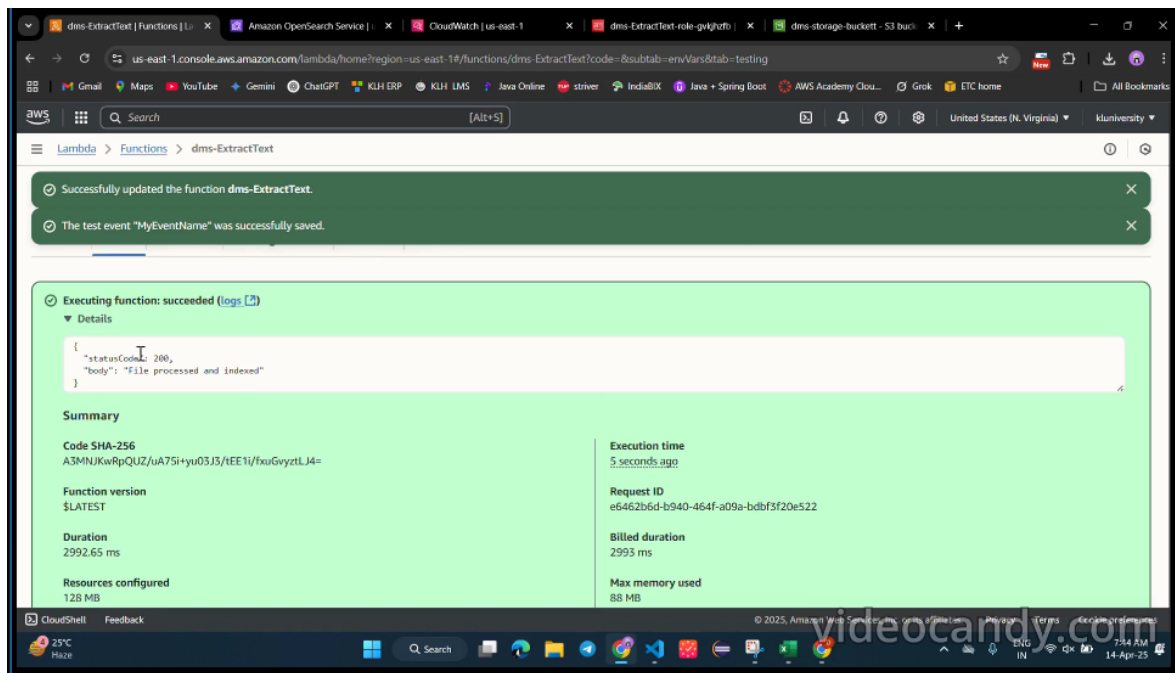


Fig 4.5.2 Checking OpenSearch index for indexed content.

Step 6 : Monitor System Performance

- Use CloudWatch to monitor Lambda executions, API Gateway requests, and OpenSearch metrics.
- Check logs for errors like NoSuchKey or NameError.

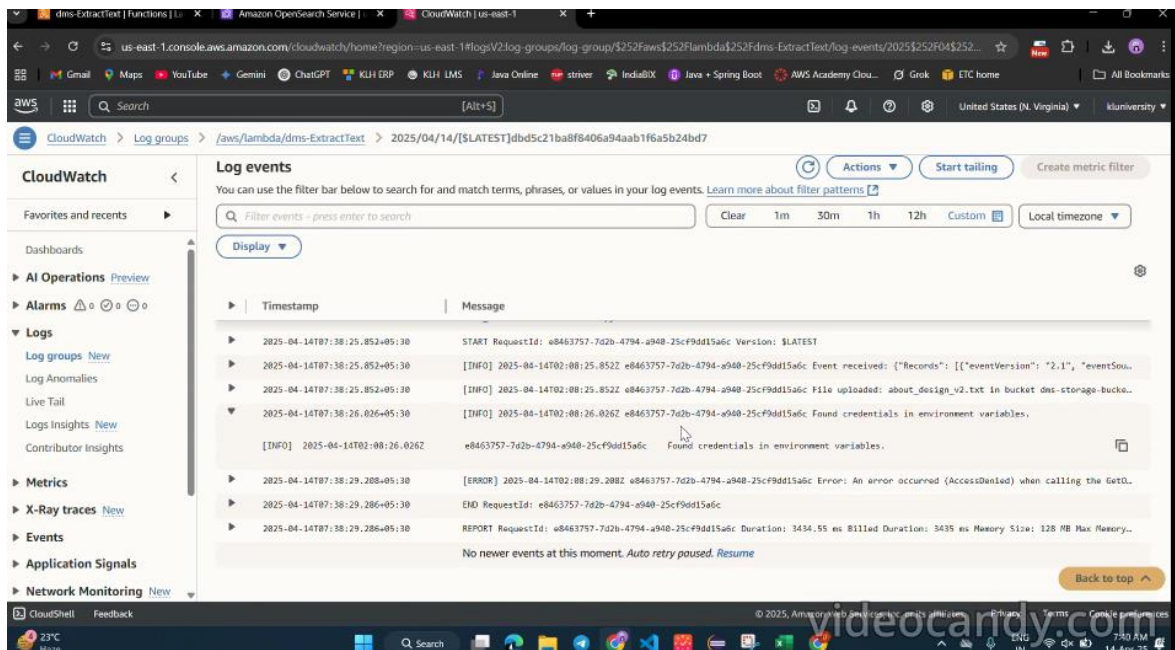


Fig 4.6 CloudWatch logs for Lambda execution.

Step 7 : Test Search Functionality

- Query the API with keywords like “design” or “emotion ” using Postman.
- Verify search results and optimize query performance.

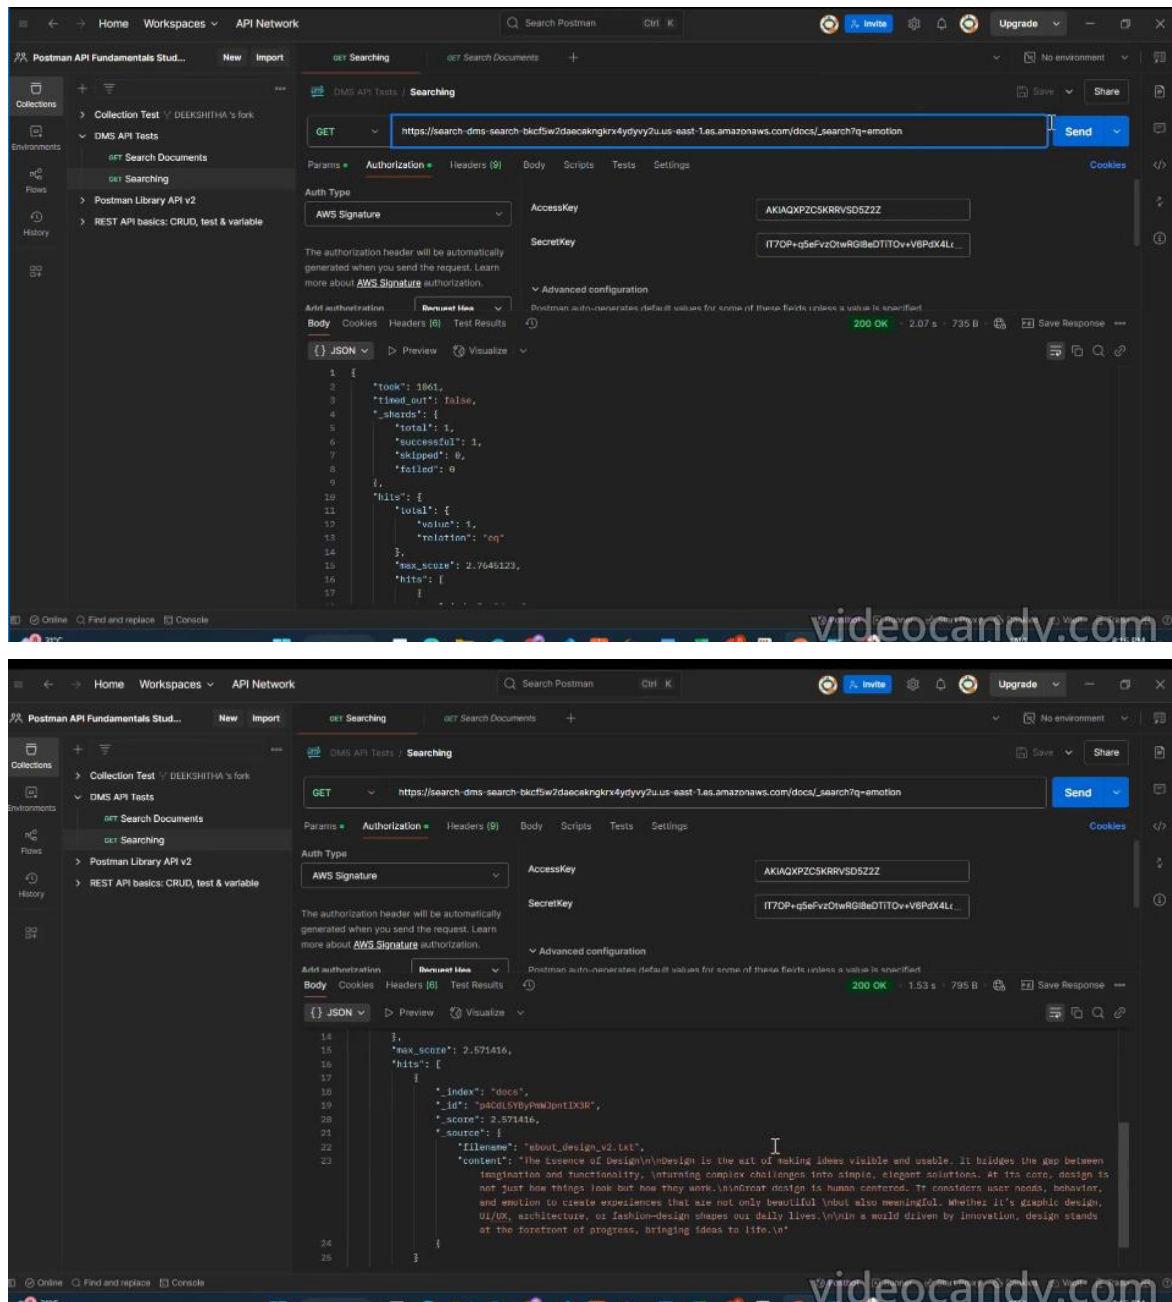


Fig 4.7.1 Postman response with search results.

5. LEARNING OUTCOMES

Through the development of the Cloud-Based Document Management System with AWS Elasticsearch, several key skills were acquired:

- Understanding AWS OpenSearch for Full-Text Search

Gained hands-on experience in indexing and searching document content using Amazon OpenSearch Service, enabling efficient retrieval of documents by keywords [1].

- Serverless Automation with AWS Lambda

Learned to automate document processing and indexing using Lambda, including text extraction with pymupdf and handling S3 events [7].

- Serverless Automation with AWS Lambda

Mastered the use of S3 for scalable document storage and IAM for secure access control, resolving issues like 403 Forbidden errors [5][8].

- API Development with AWS API Gateway

Developed skills in building RESTful APIs for search functionality, testing with Postman, and debugging issues like DNS errors [6].

- System Monitoring with Amazon CloudWatch

Enhanced proficiency in monitoring system performance and debugging errors using CloudWatch logs and metrics [5].

- Scalability and Fault Tolerance in Serverless Architectures

Designed a fault-tolerant, scalable system using serverless AWS services, optimizing for performance and cost-efficiency [5].

This project deepened the understanding of cloud-native architectures and practical applications of AWS services in document management.

6. CONCLUSION

The Cloud-Based Document Management System with AWS Elasticsearch demonstrated the power of AWS services like S3, OpenSearch, Lambda, and API Gateway in building a scalable, serverless solution for document storage, indexing, and search. The project successfully enabled the upload, indexing, and keyword-based retrieval of documents like “about_design_v2”, with a secure and automated workflow. It addressed challenges like permissions errors, dependency management, and search accuracy, showcasing the design of fault-tolerant and scalable cloud architectures. This hands-on experience strengthened skills in cloud computing, serverless automation, and search technologies, providing a solid foundation for future advancements in cloud-based document management and enterprise applications.

7. REFERENCES

- [1] Amazon OpenSearch Service Documentation:
<https://docs.aws.amazon.com/opensearch-service/latest/developerguide/what-is.html>
- [2] AWS S3 Documentation:
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html>
- [3] AWS Lambda Documentation:
<https://docs.aws.amazon.com/lambda/latest/dg/welcome.html>
- [4] AWS API Gateway Documentation:
<https://docs.aws.amazon.com/apigateway/latest/developerguide/welcome.html>
- [5] AWS CloudWatch Documentation:
<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/WhatIsCloudWatch.html>
- [6] AWS IAM Documentation:
<https://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
- [7] PyMuPDF Documentation:
<https://pymupdf.readthedocs.io/en/latest/>
- [8] Troubleshooting Amazon OpenSearch Service:
<https://docs.aws.amazon.com/opensearch-service/latest/developerguide/troubleshooting.html>
- [9] Building a Serverless Document Management System with AWS:
<https://aws.amazon.com/blogs/compute/building-a-serverless-document-management-system/>